

# Predicting Human Decision-making in Games

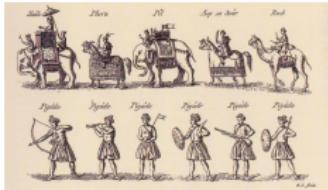
Who? Charles Zheng

From? Stanford University

When? May 2, 2016

# History of Chess

Shatranj ( $\approx$ 500 AD)



Chess (1450 AD)



Shogi ( $\approx$ 1500 AD)



# Doubutsu Shogi (Animal Shogi)



2009, Madoka Kitao





ピ ウ ブ ツ シ ょ う き  
LET'S CATCH THE LION!









ピ ウ ハ ツ シ ょ う き  
LET'S CATCH THE LION!



A

B

C



1

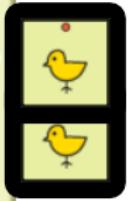
2

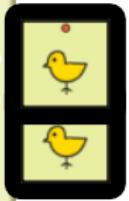
3

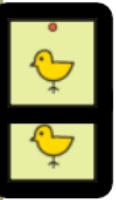
4











# What is a game?

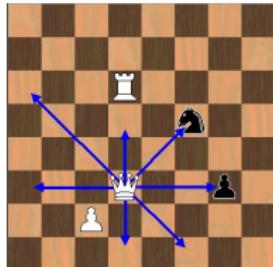
- Set of states  $s \in S$ .
- Set of winning states for player 1,  $W_1 \subset S$ .
- Winning states for player 2,  $W_2 \subset S$ .
- Each state has a set of legal actions  $\mathcal{A}(s)$ .
- Player 1 and player 2 take turns choosing the action.
- A *transition function*  $P_a(s, s')$  determines the next state  $s'$  resulting from taking action  $a$  in states  $s$ .
- In *deterministic games*, the transition function equals one for one  $s'$  and zero otherwise.

## Recorded Game Data

- List of moves made by both players: "1. Pawn from e2 to e4, 2. Pawn from e7 to e5."
- Convert to list of *state-action* pairs,  $(s_i, a_i)$ .
- $s_i$  is state of game at the beginning of turn  $i$ ,  $a_i$  is the move selected at turn  $i$ .

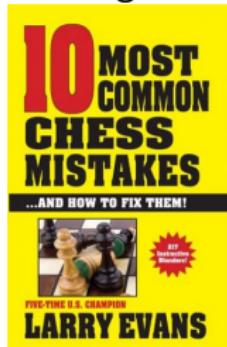
# The prediction problem

- Database of games from players  $p_1, \dots, p_m$  with:
  - Date played
  - Player 1 and player 2 identities
  - List of moves → player, state, action pairs  $(p_j, s_i, a_i)$
- In a *new game*, can we predict which move a given player  $p$  will make, given the game state  $s$ ?
- In other words: given  $(p, s)$ , guess  $a \in \mathcal{A}(s)$ .



# Motivation

- First step for building superhuman AI player
- Detecting computer-aided cheating (Regan 2013)
- Objectively evaluating professional players (Matej 2011)
- Betting on results of professional games?
- Writing a book about ‘most common chess mistakes’



- Use machine learning to learn about human learning (?)

## Section 2

### Methods

# Features

- In order to apply machine learning, need numeric features for the state  $s$ .
- Let  $x_1(s), \dots, x_p(s)$  denote the features,  $\vec{x}(s)$ =feature vector.
- Borrow features from chess programming: material, board control, king safety, etc.?
- Or try to do generic *feature selection* or *representation learning*?

## Features

We use a minimalistic featurization. The 3x4 board is converted into a 136-length binary vector.

$x_1$	Player 1 has a king at (1,1)?
$x_2$	Player 1 has a rook at (1,1)?
$x_{23}$	Player 2 has a bishop at (1, 3)?
$x_{132}$	Does player 1 have two rooks in hand?

We also consider *second-order interactions*: e.g.  $x_4x_{23} =$   
Does player 1 have a pawn on (1,1) and player 2 have a  
bishop on (1, 3)?

## Policy model

- Fit the model

$$\Pr[A = a|s] = \frac{\exp[\beta_a^T \vec{x}(s)]}{\sum_{a \in \mathcal{A}} \exp[\beta_a^T \vec{x}(s)]}.$$

where  $\mathcal{A}$  is the set of *all possible moves* in the game (not just legal moves in  $s$ ).

- When *predicting*, choose the action with the highest predicted probability among *legal* actions

$$\hat{A}(s) = \max_{a \in \mathcal{A}(s)} \beta_a^T \vec{x}(s).$$

## Policy and values

- A policy  $\pi(s, a)$  specifies a probability distribution of *actions* for each state  $s \in S$ .
- The value of a state  $V^{\pi_1, \pi_2}(s)$  is the probability that player 1 wins if player 1 uses policy  $\pi_1$  and player 2 uses policy  $\pi_2$ .

$$V^{\pi_1, \pi_2}(s) = \begin{cases} 1 & \text{if } s \in W_1 \\ 0 & \text{if } s \in W_2 \\ \sum_{a \in \mathcal{A}} \pi_i(s, a) \sum_{s'} P_a(s, s') V^{\pi_1, \pi_2}(s') & \text{otherwise} \end{cases}$$

where  $i = 1$  if it's player 1's turn and  $i = 2$  if it's player 2's turn.

## Evaluation functions

- In a deterministic game, let  $s'(s, a)$  denote the state with probability 1 resulting from  $(s, a)$
- Suppose player 1 knows player 2's policy. Then it would be rational for player 1 to choose  $a$  which maximizes the

$$V^{\pi_1, \pi_2}(s'(s, a)).$$

- However, humans are not perfectly rational nor omniscient. Perhaps players have an intuitive *evaluation function*  $E$  which approximates the true value function,

$$E(s) \approx V^{\pi_1, \pi_2}(s).$$

## Evaluation model

- Suppose a player does have a mental evaluation function  $E$ . Should they always choose  $a$  which maximizes  $E(s'(s, a))$ ?
- Even in deterministic games, there is an advantage to being unpredictable!
- This leads to a multinomial model of player choice:

$$\Pr[A = a|s] = \frac{\exp[E(s'(s, a))]}{\sum_{a' \in \mathcal{A}(s)} E(s'(s, a'))}$$

Note that  $E(s)$  need not be a probability: it could be a real number. Real-valued  $E(s)$  may be more realistic, anyways.

## How to fit the evaluation model?

- For fixed player  $p$ , let  $\{(s_i, a_i)\}_{i=1}^n$  denote all the state-action pairs for that player in the database.
- Fit a logistic evaluation model, minimizing the loss

$$\sum_{i=1}^n \vec{x}(s'(s_i, a_i))^T \beta - \log\left(\sum_{a \in \mathcal{A}} \exp[\beta^T \vec{x}(s'(s_i, a))]\right).$$

- Loss function is convex!
- For large feature models, recommended to use L1 and L2 regularization.
- To predict, choose  $a \in \mathcal{A}$  maximizing  $\beta^T s'(s, a)$ .

## Section 3

Data

# Doubutsu Shogi on LittleGolem

- Data obtained from [littlegolem.com](http://littlegolem.com), a free turn-based game site, using `scrapeR`
- 85 players, 727 games, 17094 turns (state-action pairs)
- Oct 2012 to Apr 2016
- Usernames have been anonymized

