

Predicting Human Decision-making in Games

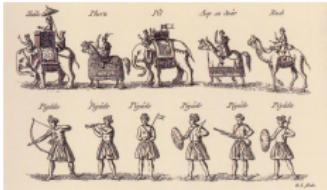
Who? Charles Zheng

From? Stanford University

When? May 2, 2016

History of Chess

Shatranj (\approx 500 AD)



Chess (1450 AD)



Shogi (\approx 1500 AD)

Doubutsu Shogi (Animal Shogi)



2009, Madoka Kitao









ピ ウ ブ ツ シ ょ う き
LET'S CATCH THE LION!







ピ ウ ハ ツ シ ょ う き
LET'S CATCH THE LION!



A

B

C

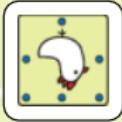


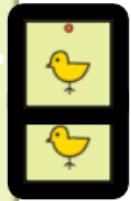
1

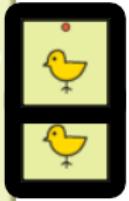
2

3

4











What is a game?

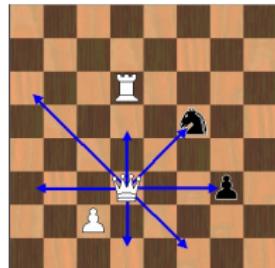
- Set of states $s \in S$.
- Set of winning states for player 1, $W_1 \subset S$.
- Winning states for player 2, $W_2 \subset S$.
- Each state has a set of legal actions $\mathcal{A}(s)$.
- Player 1 and player 2 take turns choosing the action.
- A *transition function* $P_a(s, s')$ determines the next state s' resulting from taking action a in states s .
- In *deterministic games*, the transition function equals one for one s' and zero otherwise.

Recorded Game Data

- List of moves made by both players: "1. Pawn from e2 to e4, 2. Pawn from e7 to e5."
- Convert to list of *state-action* pairs, (s_i, a_i) .
- s_i is state of game at the beginning of turn i , a_i is the move selected at turn i .

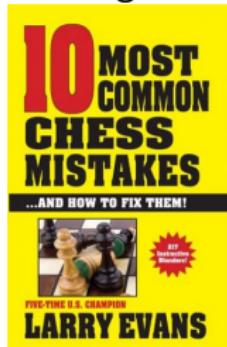
The prediction problem

- Database of games from players p_1, \dots, p_m with:
 - Date played
 - Player 1 and player 2 identities
 - List of moves → player, state, action pairs (p_j, s_i, a_i)
- In a *new game*, can we predict which move a given player p will make, given the game state s ?
- In other words: given (p, s) , guess $a \in \mathcal{A}(s)$.



Motivation

- First step for building superhuman AI player
- Detecting computer-aided cheating (Regan 2013)
- Objectively evaluating professional players (Matej 2011)
- Betting on results of professional games?
- Writing a book about ‘most common chess mistakes’



- Use machine learning to learn about human learning (?)

Section 2

Methods

Features

- In order to apply machine learning, need numeric features for the state s .
- Let $x_1(s), \dots, x_p(s)$ denote the features, $\vec{x}(s)$ =feature vector.
- Borrow features from chess programming: material, board control, king safety, etc.?
- Or try to do generic *feature selection* or *representation learning*?

Features

We use a minimalistic featurization. The 3x4 board is converted into a 136-length binary vector.

x_1	Player 1 has a king at (1,1)?
x_2	Player 1 has a rook at (1,1)?
x_{23}	Player 2 has a bishop at (1, 3)?
x_{132}	Does player 1 have two rooks in hand?

We also consider *second-order interactions*: e.g. $x_4x_{23} =$
Does player 1 have a pawn on (1,1) and player 2 have a
bishop on (1, 3)?

Policy model

- Fit the model

$$\Pr[A = a|s] = \frac{\exp[\beta_a^T \vec{x}(s)]}{\sum_{a \in \mathcal{A}} \exp[\beta_a^T \vec{x}(s)]}.$$

where \mathcal{A} is the set of *all possible moves* in the game (not just legal moves in s).

- When *predicting*, choose the action with the highest predicted probability among *legal* actions

$$\hat{A}(s) = \max_{a \in \mathcal{A}(s)} \beta_a^T \vec{x}(s).$$

Policy and values

- A policy $\pi(s, a)$ specifies a probability distribution of *actions* for each state $s \in S$.
- The value of a state $V^{\pi_1, \pi_2}(s)$ is the probability that player 1 wins if player 1 uses policy π_1 and player 2 uses policy π_2 .

$$V^{\pi_1, \pi_2}(s) = \begin{cases} 1 & \text{if } s \in W_1 \\ 0 & \text{if } s \in W_2 \\ \sum_{a \in \mathcal{A}} \pi_i(s, a) \sum_{s'} P_a(s, s') V^{\pi_1, \pi_2}(s') & \text{otherwise} \end{cases}$$

where $i = 1$ if it's player 1's turn and $i = 2$ if it's player 2's turn.

Evaluation functions

- In a deterministic game, let $s'(s, a)$ denote the state with probability 1 resulting from (s, a)
- Suppose player 1 knows player 2's policy. Then it would be rational for player 1 to choose a which maximizes the

$$V^{\pi_1, \pi_2}(s'(s, a)).$$

- However, humans are not perfectly rational nor omniscient. Perhaps players have an intuitive *evaluation function* E which approximates the true value function,

$$E(s) \approx V^{\pi_1, \pi_2}(s).$$

- This is also how chess engines work—idea first proposed by Shannon (1949)

Evaluation model

- Suppose a player does have a mental evaluation function E . Should they always choose a which maximizes $E(s'(s, a))$?
- Even in perfect information games, there is an advantage to being unpredictable!
- This leads to a multinomial model of player choice:

$$\Pr[A = a|s] = \frac{\exp[E(s'(s, a))]}{\sum_{a' \in \mathcal{A}(s)} E(s'(s, a'))}$$

Note that $E(s)$ need not be a probability: it could be a real number. Real-valued $E(s)$ may be more realistic, anyways.

How to fit the evaluation model?

- For fixed player p , let $\{(s_i, a_i)\}_{i=1}^n$ denote all the state-action pairs for that player in the database.
- Fit a logistic evaluation model, minimizing the loss

$$-\sum_{i=1}^n \vec{x}(s'(s_i, a_i))^T \beta - \log\left(\sum_{a \in \mathcal{A}} \exp[\beta^T \vec{x}(s'(s_i, a))]\right).$$

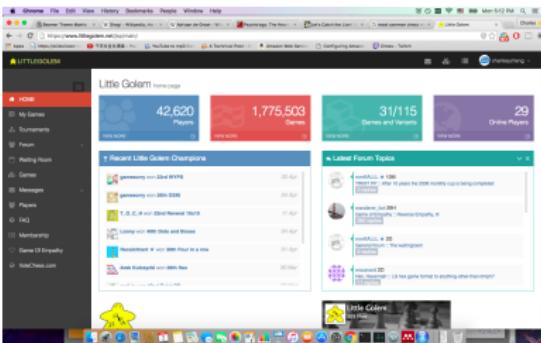
- Loss function is convex!
- For large feature models, recommended to use L1 and L2 regularization.
- To predict, choose $a \in \mathcal{A}$ maximizing $\beta^T s'(s, a)$.

Section 3

Data

Doubutsu Shogi on LittleGolem

- Data obtained from littlegolem.com, a free turn-based game site, using `scrapeR`
- 85 players, 727 games, 17094 turns (state-action pairs)
- Oct 2012 to Apr 2016
- Usernames have been anonymized



Ranking the players

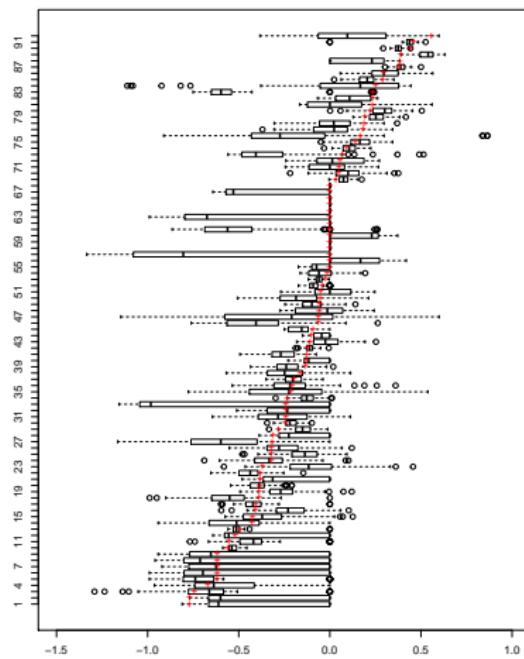
- θ_i parameterizes the skill level of player i

$$\Pr[i \text{ wins} | i \text{ vs } j] = \frac{\exp[\theta_i - \theta_j]}{1 + \exp[\theta_i - \theta_j]}.$$

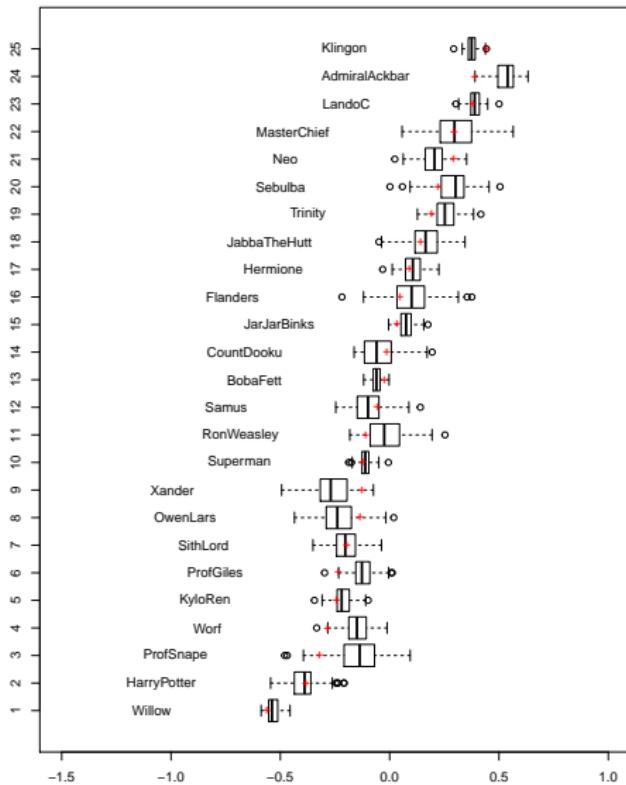
Bradley-Terry model

- A spread of $\theta_i - \theta_j = 1$ means a 73% win rate for i .
- Fit using `glmnet`: add some regularization
- Bootstrap to get ‘confidence’ intervals

All players



Players with > 5 games.



How close are we to perfect play?

The game is a theoretical loss for player 1. We looked at games where both players are above a given skill threshold... does the win % approach the theoretical value?

Group	# of games	% win for P1
All	701	0.48
$\theta > 0$	119	0.44
$\theta > 0.2, > 5$ games	19	0.31
$\theta > 0.3, > 5$ games	10	0.40
Theoretical		0

Move prediction

First step: lump all of the players into one group. How well can we predict?

Method	No. params	Accuracy
Markov model	186	0.441
Policy, order-1	25k	0.565
Policy, order-2	1.7m	0.559
Eval, order-1	136	0.370
Eval, order-2	9316	0.559
Eval, order-3	419k	0.439

Note: regularization was done in an ad hoc way.

Move prediction

What happens if we combine the policy and evaluation models?

$$\hat{\Pr}[a|s] = 0.5\hat{\Pr}_{policy}[a|s] + 0.5\hat{\Pr}_{eval}[a|s]$$

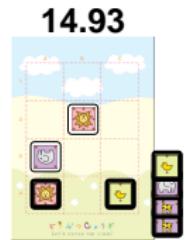
Method	No. params	Accuracy
Policy, order-1	25k	0.565
Eval, order-2	9316	0.559
Ensemble	34k	0.611

Validating the evaluation function

- The learned eval function should estimate some monotonic function of the probability of a win.
- Take game states near the ends of games:
 - if player 1 won: $E(s)$ should be high.
 - if player 1 lost: $E(s)$ should be low.
- This gives a somewhat independent validation, since the win vs lose information was *not* used to fit the evaluation model!
- We apply this for the second-order model (9316 params, 0.559 accuracy)

Validating the evaluation function

Win

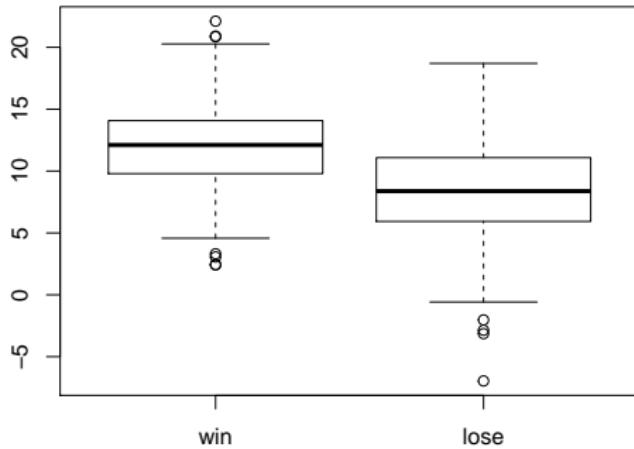


Lose



Validating the evaluation function

Game states 3 or 4 moves from the end of the game, on Player 1's move.



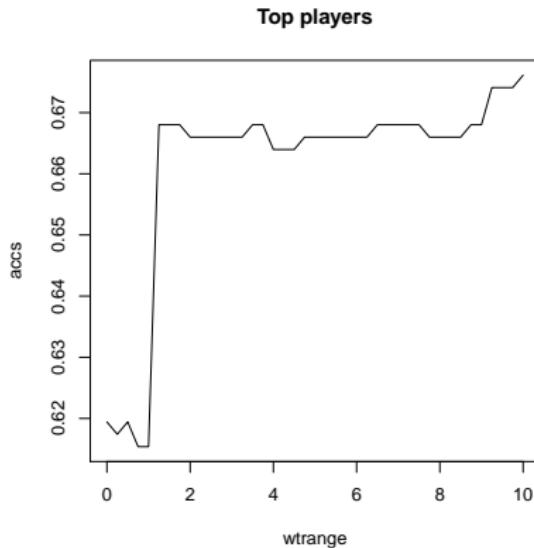
Personalized Prediction

- Data (p_i, s_i, a_i) where p_i is the player
- Let P be a set of 'similar' players.
- We would like to predict moves made by players in P .
- Fit the evaluation model by minimizing

$$-\sum_{i=1}^n w_i (\vec{x}(s'(s_i, a_i))^T \beta - \log(\sum_{a \in \mathcal{A}} \exp[\beta^T \vec{x}(s'(s_i, a))])).$$

- Set $w_i = 1 + w$ if $p \in P$ and $w_i = 1$ for $p \notin P$. (Then normalize so $\sum w_i = n$).
- w controls the extra weight given to players in P .

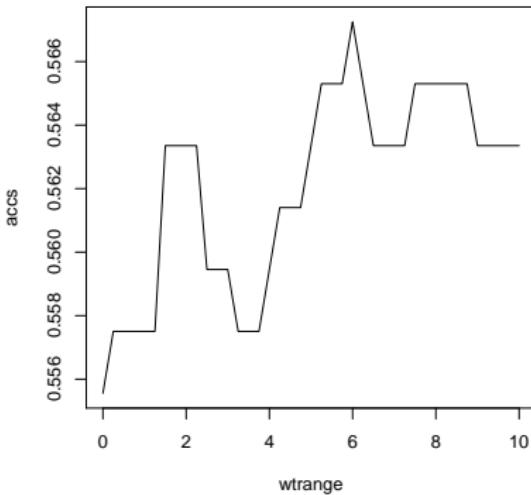
Personalized Prediction



Players with $\theta > 0.2$ and more than 5 games (6 players).

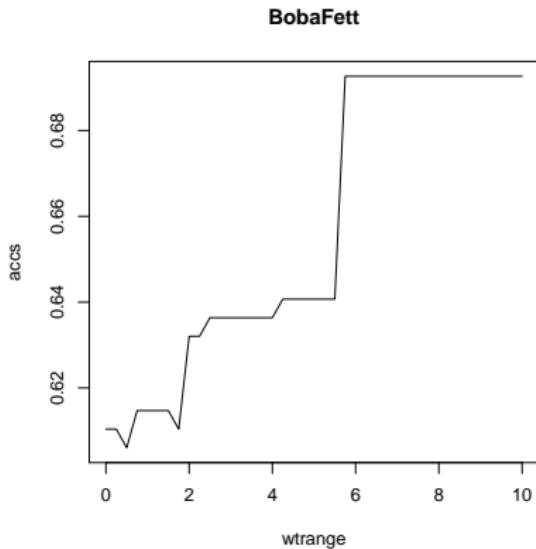
Personalized Prediction

Least-rated players



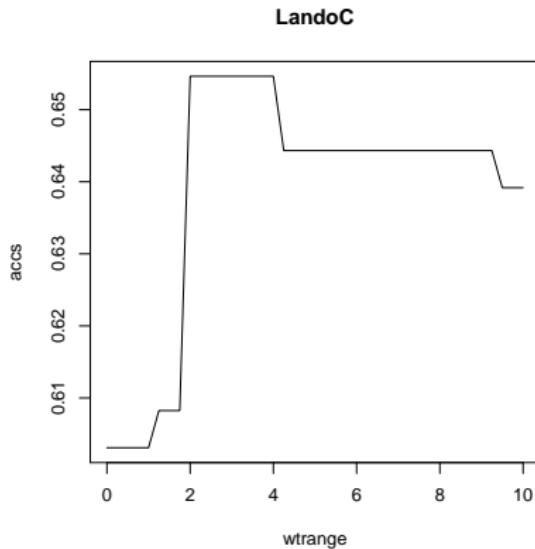
Players with $\theta < 0.2$ and more than 5 games (16 players).

Personalized Prediction



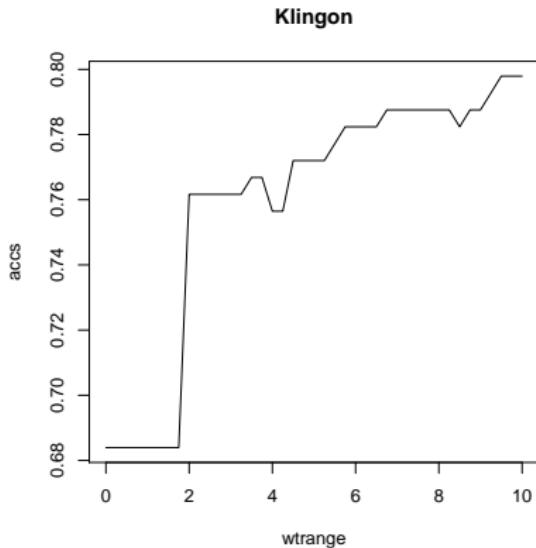
BobaFett, 78 wins - 118 losses, $\theta = -0.02$.

Personalized Prediction



LandoC, 108 wins - 7 losses, $\theta = 0.38$.

Personalized Prediction



Klingon, 98 wins - 7 losses, $\theta = 0.44$.

Conclusions

- Two approaches for prediction problem: policy model and evaluation model.
- Policy model works well for the most part, but may fare worse in new situations.
- Evaluation model may be more parameter-efficient.
- Ensemble works better than either model.
- Second-order evaluation model could likely be improved, based on validation results.
- Hypothesis: better players are more predictable than worse players?

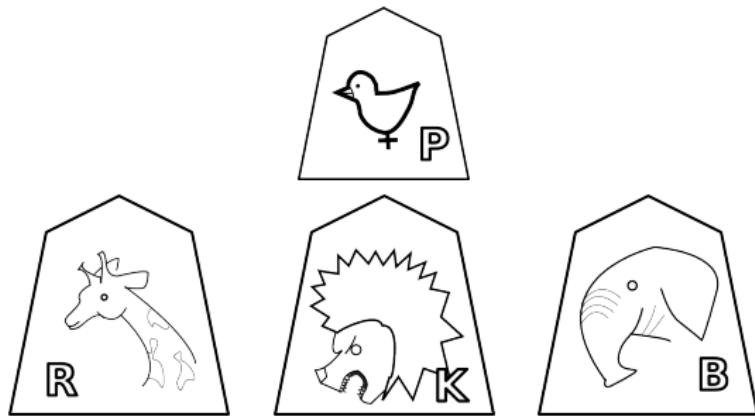
Future play

As opposed to future ‘work’.

- More games! More complex games and different types of games!
- Decision trees might be very, very good!
- ‘Regularize’ the evaluation functions using the assumption of self-consistency

$$E(s) \approx \min_{a \in \mathcal{A}(s)} \max_{a' \in \mathcal{A}(s')} E(s'(s', a')).$$

- Track player improvement over time.
- Our predictive models should be better than using a computer player to predict human moves, right?



Thanks!